

Restaurants receive thousands of customer reviews, but star ratings alone fail to explain why ratings increase or decline. Reviews often contain rich information about food quality, service, pricing, ambience, and operational factors, yet this information remains unstructured and difficult to analyze at scale. This limits restaurants' ability to identify the drivers of customer satisfaction and dissatisfaction.

Notebook 3: Insight Analysis

This notebook performs exploratory and diagnostic analysis on a previously constructed review-level dataset (aspect_df) to explain why restaurant ratings vary.

By analyzing sentiment distributions across review topics (aspects) and linking them to restaurant operational attributes, this notebook identifies the key drivers of positive and negative customer sentiment. The goal is to translate unstructured review text into actionable insights that help explain changes in star ratings beyond the ratings themselves.

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

Load Prepared Analytical Dataset

```
import pandas as pd
import pickle

PATH = "/content/drive/MyDrive/restaurant_sentiment_model/aspect_df.pkl"

#load
aspect_df = pd.read_pickle(PATH)
```

aspect_df.head()

	business_id	review	topic	sentiment	AcceptsInsurance	AgesAllowed	Alcohol	Ambience	BYOB
0	YtSqYv1Q_pOltsVPSx54SA	Tremendous service (Big shout out to Douglas) ...		-1	2	Unknown	Unknown	'full_bar'	{'romantic': False, 'intimate': False, 'classy': ...}
1	aY_n9RSaD2Yw09jSFFePew	We visited once and were very disappointed in ...		3	0	Unknown	Unknown	u'beer_and_wine'	{'divey': False, 'hipster': False, u'casual': ...}
2	18eWJFJbXyR9j_5xfcRLYA	This is the first time I tried this place and ...		6	2	Unknown	Unknown	u'beer_and_wine'	{'divey': False, 'hipster': False, u'casual': ...}
3	jOOOrH5n2ijnsZKxzPSAiw	This is one of the busiest Chick fil A's I've ...		-1	2	Unknown	Unknown	u'none'	{'romantic': False, 'intimate': False, 'touris...}
4	1QVB0_-piu0GXes87BXeGw	Love this place...best hot dogs and chili dogs...		-1	2	Unknown	Unknown	u'none'	{'touristy': False, 'hipster': False, 'romanti...}

5 rows × 43 columns

aspect_df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 43 columns):
 #   Column           Non-Null Count Dtype  
 --- 
 0   business_id      10000 non-null  object  
 1   review           10000 non-null  object  
 2   topic            10000 non-null  int64  
 3   sentiment        10000 non-null  int64
```

```

4   AcceptsInsurance      10000 non-null object
5   AgesAllowed          10000 non-null object
6   Alcohol              10000 non-null object
7   Ambience             10000 non-null object
8   BYOB                 10000 non-null object
9   BYOBCorkage          10000 non-null object
10  BestNights           10000 non-null object
11  BikeParking          10000 non-null object
12  BusinessAcceptsBitcoin 10000 non-null object
13  BusinessAcceptsCreditCards 10000 non-null object
14  BusinessParking       10000 non-null object
15  ByAppointmentOnly    10000 non-null object
16  Caters               10000 non-null object
17  CoatCheck             10000 non-null object
18  Corkage              10000 non-null object
19  DietaryRestrictions  10000 non-null object
20  DogsAllowed          10000 non-null object
21  DriveThru            10000 non-null object
22  GoodForDancing        10000 non-null object
23  GoodForKids           10000 non-null object
24  GoodForMeal           10000 non-null object
25  HairSpecializesIn    10000 non-null object
26  HappyHour             10000 non-null object
27  HasTV                10000 non-null object
28  Music                 10000 non-null object
29  NoiseLevel            10000 non-null object
30  Open24Hours           10000 non-null object
31  OutdoorSeating        10000 non-null object
32  RestaurantsAttire     10000 non-null object
33  RestaurantsCounterService 10000 non-null object
34  RestaurantsDelivery   10000 non-null object
35  RestaurantsGoodForGroups 10000 non-null object
36  RestaurantsPriceRange2 10000 non-null object
37  RestaurantsReservations 10000 non-null object
38  RestaurantsTableService 10000 non-null object
39  RestaurantsTakeOut    10000 non-null object
40  Smoking               10000 non-null object
41  WheelchairAccessible  10000 non-null object
42  WiFi                  10000 non-null object
dtypes: int64(2), object(41)
memory usage: 3.3+ MB

```

```
pip install -U bertopic
```

```

Requirement already satisfied: scikit-learn>=1.0 in /usr/local/lib/python3.12/dist-packages (from bertopic) (1.6.1)
Requirement already satisfied: sentence-transformers>=0.4.1 in /usr/local/lib/python3.12/dist-packages (from bertopic) (5.
Requirement already satisfied: tqdm>=4.41.1 in /usr/local/lib/python3.12/dist-packages (from bertopic) (4.67.1)
Requirement already satisfied: llvmlite>0.36.0 in /usr/local/lib/python3.12/dist-packages (from bertopic) (0.43.0)
Requirement already satisfied: scipy>=1.0 in /usr/local/lib/python3.12/dist-packages (from hdbscan>=0.8.29->bertopic) (1.1
Requirement already satisfied: joblib>=1.0 in /usr/local/lib/python3.12/dist-packages (from hdbscan>=0.8.29->bertopic) (1.
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.12/dist-packages (from pandas>=1.1.5->bert
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-packages (from pandas>=1.1.5->bertopic) (202
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas>=1.1.5->bertopic) (2
Requirement already satisfied: tenacity>=6.2.0 in /usr/local/lib/python3.12/dist-packages (from plotly>=4.7.0->bertopic) (
Requirement already satisfied: packaging in /usr/local/lib/python3.12/dist-packages (from plotly>=4.7.0->bertopic) (25.0)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn>=1.0->be
Requirement already satisfied: transformers<6.0.0,>=4.1.0 in /usr/local/lib/python3.12/dist-packages (from sentence-trans
Requirement already satisfied: torch>=1.11.0 in /usr/local/lib/python3.12/dist-packages (from sentence-transformers)>=0.4.1
Requirement already satisfied: huggingface-hub>=0.20.0 in /usr/local/lib/python3.12/dist-packages (from sentence-transform
Requirement already satisfied: typing_extensions>=4.5.0 in /usr/local/lib/python3.12/dist-packages (from sentence-transfor
Requirement already satisfied: numba>=0.51.2 in /usr/local/lib/python3.12/dist-packages (from umap-learn>=0.5.0->bertopic)
Requirement already satisfied: pynndescent>=0.5 in /usr/local/lib/python3.12/dist-packages (from umap-learn>=0.5.0->bertop
Requirement already satisfied: filelock in /usr/local/lib/python3.12/dist-packages (from huggingface-hub>=0.20.0->sentence
Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.12/dist-packages (from huggingface-hub>=0.20.0->
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.12/dist-packages (from huggingface-hub>=0.20.0->sente
Requirement already satisfied: requests in /usr/local/lib/python3.12/dist-packages (from huggingface-hub>=0.20.0->sentence
Requirement already satisfied: hf-xet<2.0.0,>=1.1.3 in /usr/local/lib/python3.12/dist-packages (from huggingface-hub>=0.20
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.8.2->pandas>=1
Requirement already satisfied: setuptools in /usr/local/lib/python3.12/dist-packages (from torch>=1.11.0->sentence-transfo
Requirement already satisfied: sympy>=1.13.3 in /usr/local/lib/python3.12/dist-packages (from torch>=1.11.0->sentence-tran
Requirement already satisfied: networkx>=2.5.1 in /usr/local/lib/python3.12/dist-packages (from torch>=1.11.0->sentence-tr
Requirement already satisfied: jinja2 in /usr/local/lib/python3.12/dist-packages (from torch>=1.11.0->sentence-transformer
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages (from torch>=1.1
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages (from torch>=1.1
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.6.80 in /usr/local/lib/python3.12/dist-packages (from torch>=1.1
Requirement already satisfied: nvidia-cudnn-cu12==9.10.2.21 in /usr/local/lib/python3.12/dist-packages (from torch>=1.11.0
Requirement already satisfied: nvidia-cUBLAS-cu12==12.6.4.1 in /usr/local/lib/python3.12/dist-packages (from torch>=1.11.0
Requirement already satisfied: nvidia-cufft-cu12==11.3.0.4 in /usr/local/lib/python3.12/dist-packages (from torch>=1.11.0-
Requirement already satisfied: nvidia-curand-cu12==10.3.7.77 in /usr/local/lib/python3.12/dist-packages (from torch>=1.11.
Requirement already satisfied: nvidia-cusolver-cu12==11.7.1.2 in /usr/local/lib/python3.12/dist-packages (from torch>=1.11
Requirement already satisfied: nvidia-cusparse-cu12==12.5.4.2 in /usr/local/lib/python3.12/dist-packages (from torch>=1.11
Requirement already satisfied: nvidia-cusparseL-cu12==0.7.1 in /usr/local/lib/python3.12/dist-packages (from torch>=1.11.

```

Installing collected packages: bertopic
154.7/154.7 KB 8.2 MB/s eta 0:00:00

```
Successfully installed bertopic-0.17.4
```

```
from bertopic import BERTopic
from sentence_transformers import SentenceTransformer

embedding_model = SentenceTransformer("all-MiniLM-L6-v2")
MODEL_PATH = "/content/drive/MyDrive/restaurant_sentiment_model/aspect"
topic_model = BERTopic.load(
    MODEL_PATH,
    embedding_model=embedding_model
)

/usr/local/lib/python3.12/dist-packages/hdbscan/robust_single_linkage_.py:175: SyntaxWarning: invalid
    $max \{ core_k(a), core_k(b), 1/\alpha d(a,b) \}$. 349/349 [00:00<00:00, 36.0kB/s]

modules.json: 100% 116/116 [00:00<00:00, 13.8kB/s]

config_sentence_transformers.json: 100% 349/349 [00:00<00:00, 36.0kB/s]

README.md:      10.5k/? [00:00<00:00, 1.09MB/s] 53.0/53.0 [00:00<00:00, 6.39kB/s]

sentence_bert_config.json: 100% 612/612 [00:00<00:00, 63.0kB/s]

config.json: 100% 90.9M/90.9M [00:00<00:00, 163MB/s]

model.safetensors: 100% 350/350 [00:00<00:00, 37.5kB/s]

tokenizer_config.json: 100% 232k/? [00:00<00:00, 12.4MB/s] 112/112 [00:00<00:00, 13.3kB/s]

vocab.txt:      466k/? [00:00<00:00, 29.1MB/s] 190/190 [00:00<00:00, 23.7kB/s]

special_tokens_map.json: 100% 112/112 [00:00<00:00, 13.3kB/s]

config.json: 100%
```

Topic and Sentiment Label Mapping

```
import matplotlib.pyplot as plt
import pandas as pd

#Map sentiment labels
sentiment_mapping = {0: "Negative", 1: "Neutral", 2: "Positive"}
aspect_df['sentiment_label'] = aspect_df['sentiment'].map(sentiment_mapping)

#Map topic IDs to descriptive labels
topic_info = topic_model.get_topic_info()

topic_mapping = {}
for topic_id in topic_info['Topic']:
    if topic_id == -1:
        topic_mapping[topic_id] = "Other / Outliers"
    else:
        words = [word for word, _ in topic_model.get_topic(topic_id)[:3]] # top 3 words
        topic_mapping[topic_id] = " ".join(words)

aspect_df['topic_label'] = aspect_df['topic'].map(topic_mapping)
```

```
aspect_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 45 columns):
 #   Column           Non-Null Count Dtype  
 --- 
 0   business_id      10000 non-null  object  
 1   review            10000 non-null  object  
 2   topic              10000 non-null  int64   
 3   sentiment          10000 non-null  int64   
 4   AcceptsInsurance  10000 non-null  object  
 5   AgesAllowed        10000 non-null  object 
```

```

6 Alcohol          10000 non-null object
7 Ambience         10000 non-null object
8 BYOB             10000 non-null object
9 BYOBCorkage     10000 non-null object
10 BestNights      10000 non-null object
11 BikeParking     10000 non-null object
12 BusinessAcceptsBitcoin 10000 non-null object
13 BusinessAcceptsCreditCards 10000 non-null object
14 BusinessParking 10000 non-null object
15 ByAppointmentOnly 10000 non-null object
16 Caters           10000 non-null object
17 CoatCheck        10000 non-null object
18 Corkage          10000 non-null object
19 DietaryRestrictions 10000 non-null object
20 DogsAllowed      10000 non-null object
21 DriveThru        10000 non-null object
22 GoodForDancing   10000 non-null object
23 GoodForKids      10000 non-null object
24 GoodForMeal       10000 non-null object
25 HairSpecializesIn 10000 non-null object
26 HappyHour        10000 non-null object
27 HasTV            10000 non-null object
28 Music             10000 non-null object
29 NoiseLevel        10000 non-null object
30 Open24Hours       10000 non-null object
31 OutdoorSeating    10000 non-null object
32 RestaurantsAttire 10000 non-null object
33 RestaurantsCounterService 10000 non-null object
34 RestaurantsDelivery 10000 non-null object
35 RestaurantsGoodForGroups 10000 non-null object
36 RestaurantsPriceRange2 10000 non-null object
37 RestaurantsReservations 10000 non-null object
38 RestaurantsTableService 10000 non-null object
39 RestaurantsTakeOut   10000 non-null object
40 Smoking           10000 non-null object
41 WheelchairAccessible 10000 non-null object
42 WiFi              10000 non-null object
43 sentiment_label   10000 non-null object
44 topic_label        10000 non-null object
dtypes: int64(2), object(43)
memory usage: 3.4+ MB

```

```
aspect_df.head()
```

	business_id	review	topic	sentiment	AcceptsInsurance	AgesAllowed	Alcohol	Ambience	BYOB
0	YtSqYv1Q_pOltsVPSx54SA	Tremendous service (Big shout out to Douglas) ...		-1	2	Unknown	Unknown	'full_bar'	{'romantic': False, 'intimate': False, 'classy': ...}
1	aY_n9RSaD2Yw09jSFFePew	We visited once and were very disappointed in ...		3	0	Unknown	Unknown	u'beer_and_wine'	{'divey': False, 'hipster': False, 'casual': ...}
2	18eWJFJbXyR9j_5xfcRLYA	This is the first time I tried this place and ...		6	2	Unknown	Unknown	u'beer_and_wine'	{'divey': False, 'hipster': False, 'casual': ...}
3	jOOOrH5n2ijnsZKxzPSAiw	This is one of the busiest Chick fil A's I've ...		-1	2	Unknown	Unknown	u'none'	{'romantic': False, 'intimate': False, 'tourist': ...}
4	1QVB0_-piu0GXes87BXeGw	Love this place...best hot dogs and chili dogs...		-1	2	Unknown	Unknown	u'none'	{'touristy': False, 'hipster': False, 'romantic': ...}

5 rows × 45 columns

```
filtered_df = aspect_df[aspect_df["topic"] != -1]
```

```
filtered_df.head()
```

	business_id	review	topic	sentiment	AcceptsInsurance	AgesAllowed	Alcohol	Ambience	BY...
1	aY_n9RSaD2Yw09jSFFePew	We visited once and were very disappointed in ...		3	0	Unknown	Unknown	u'beer_and_wine'	{u'divey': False, u'hipster': False, u'casual': ...}
2	18eWJFJbXyR9j_5xfcRLYA	This is the first time I tried this place and ...		6	2	Unknown	Unknown	u'beer_and_wine'	{u'divey': False, u'hipster': False, u'casual': ...}
7	9OG5YkX1g2GReZM0AskizA	Great bar Happy Hour 4-7 every day. Wine & Dra...		1	2	Unknown	Unknown	'full_bar'	{u'divey': False, u'hipster': False, u'casual': ...}
8	LnQRfj3pPz0369stRnwUWw	Very good sushi. The peanut avocado roll is a ...		4	2	Unknown	Unknown	u'beer_and_wine'	{u'divey': False, u'hipster': False, u'casual': ...}
12	Mha77MN_BYGB-w7DONZ5YA	Our group of 6 decided to try something differ...		17	2	Unknown	Unknown	u'none'	{'romantic': False, 'intimate': False, 'touris...}

5 rows × 45 columns

```
import pickle

PATH = "/content/drive/MyDrive/restaurant_sentiment_model/filtered_df.pkl"

#Save
filtered_df.to_pickle(PATH)

#Later, load
filtered_df = pd.read_pickle(PATH)
```

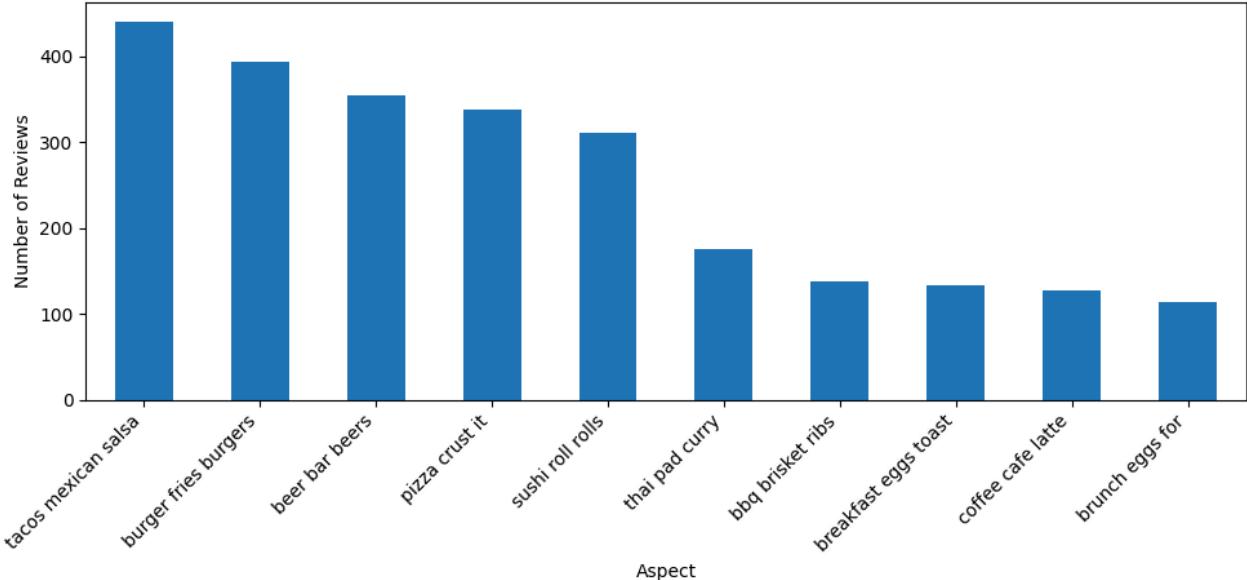
Most Discussed Restaurant Aspects

```
import matplotlib.pyplot as plt

topic_counts = (
    filtered_df["topic_label"]
    .value_counts()
    .head(10)
)

plt.figure(figsize=(10,5))
topic_counts.plot(kind="bar")
plt.title("Top 10 Most Discussed Restaurant Aspects (Excluding Outliers)")
plt.xlabel("Aspect")
plt.ylabel("Number of Reviews")
plt.xticks(rotation=45, ha="right")
plt.tight_layout()
plt.show()
```

Top 10 Most Discussed Restaurant Aspects (Excluding Outliers)



Sentiment Distribution by Restaurant Aspect

```

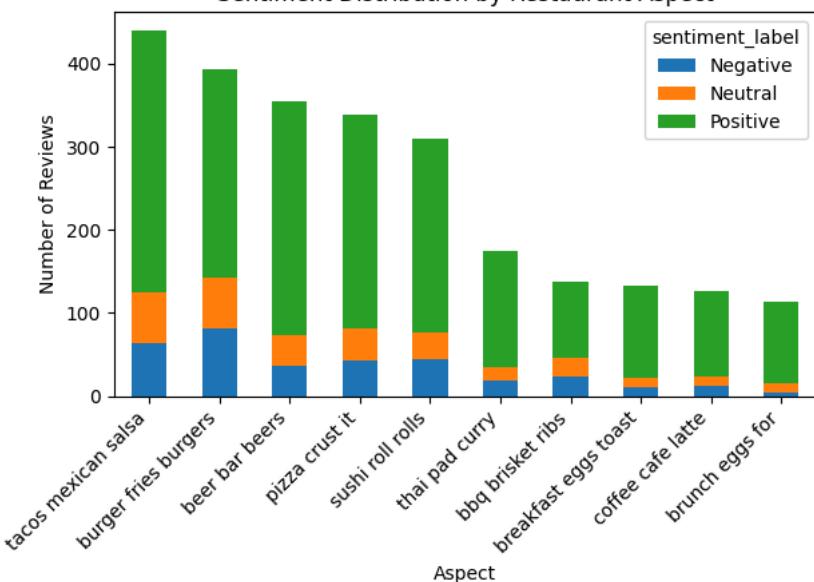
sentiment_topic = (
    aspect_df
    .groupby(["topic_label", "sentiment_label"])
    .size()
    .unstack(fill_value=0)
    .loc[topic_counts.index]
)

plt.figure(figsize=(10,5))
sentiment_topic.plot(kind="bar", stacked=True)
plt.title("Sentiment Distribution by Restaurant Aspect")
plt.xlabel("Aspect")
plt.ylabel("Number of Reviews")
plt.xticks(rotation=45, ha="right")
plt.tight_layout()
plt.show()

```

<Figure size 1000x500 with 0 Axes>

Sentiment Distribution by Restaurant Aspect



Net Sentiment of Restaurant topics

```

import matplotlib.pyplot as plt

#Get top 10 topics

```

```

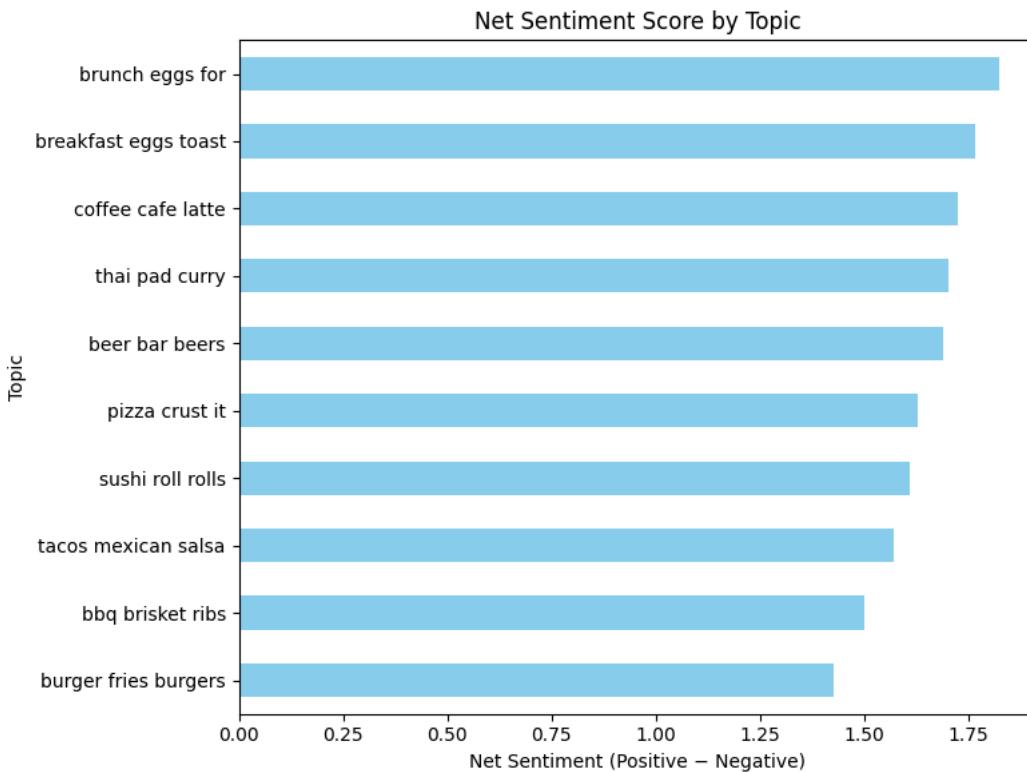
top_topics = filtered_df['topic_label'].value_counts().head(10).index

#Filter datafram to only include top topics
filtered_top_df = filtered_df[filtered_df['topic_label'].isin(top_topics)]

#Calculate mean net sentiment by topic
net_sentiment = (
    filtered_top_df
    .groupby("topic_label")["sentiment"]
    .mean()
    .sort_values()
)

#Plot horizontal bar chart
plt.figure(figsize=(8,6))
net_sentiment.plot(kind="barh", color="skyblue")
plt.title("Net Sentiment Score by Topic")
plt.xlabel("Net Sentiment (Positive - Negative)")
plt.ylabel("Topic")
plt.tight_layout()
plt.show()

```



Deep-Dive Analysis: Aspect-Specific Drivers

```

import matplotlib.pyplot as plt

# Define topic
topic = "burger fries burgers"

# Filter for topic
topic_df = filtered_df[filtered_df["topic_label"] == topic]

# Automatically detect operational attributes
exclude_cols = ["business_id", "review", "topic", "topic_label", "sentiment", "sentiment_label"]
operational_attributes = [col for col in topic_df.columns if col not in exclude_cols]

# Compute proportion of negative reviews for each attribute relative to all reviews
neg_attr_scores = {}

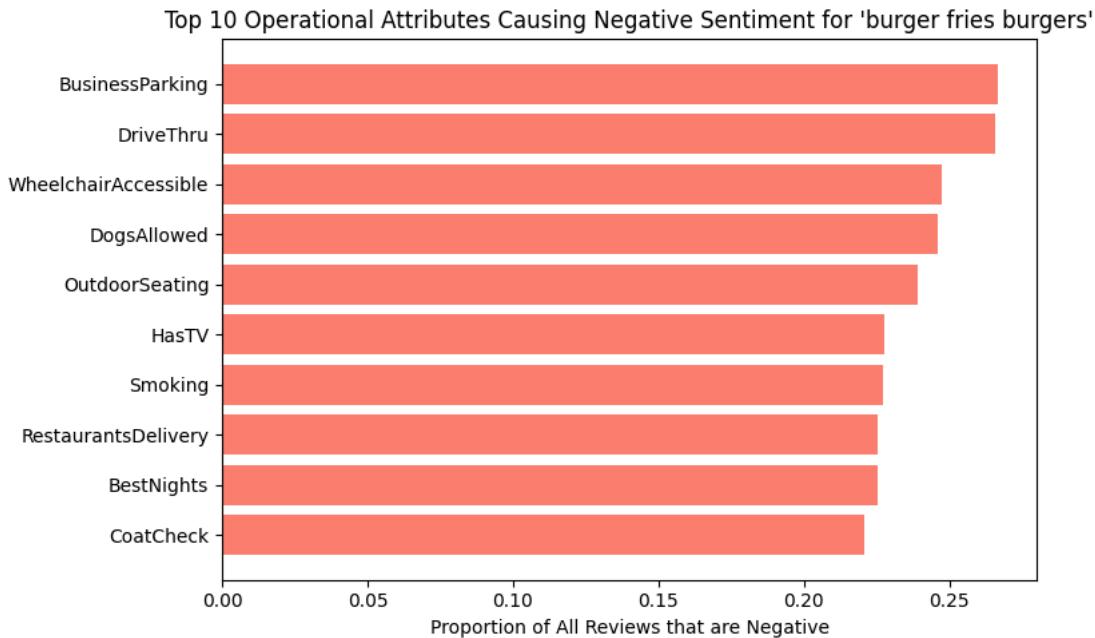
for attr in operational_attributes:
    if topic_df[attr].dtype == 'object':
        # Get all reviews for this attribute value
        counts = topic_df.groupby(attr)["sentiment_label"].value_counts(normalize=True)
        # Take the proportion of negative sentiment for the most common attribute value
        most_common_value = topic_df[attr].mode()[0]
        neg_score = counts.get((most_common_value, "Negative"), 0)
        neg_attr_scores[attr] = neg_score

```

```
# Sort descending and take top 5
top5_neg_attrs = sorted(neg_attr_scores.items(), key=lambda x: x[1], reverse=True)[:10]

# Plot
attrs, scores = zip(*top5_neg_attrs)

plt.figure(figsize=(8,5))
plt.barh(attrs, scores, color='salmon')
plt.xlabel("Proportion of All Reviews that are Negative")
plt.title(f"Top 10 Operational Attributes Causing Negative Sentiment for '{topic}'")
plt.gca().invert_yaxis() # highest at top
plt.tight_layout()
plt.show()
```



```
import matplotlib.pyplot as plt

# Define topic
topic = "burger fries burgers"

# Filter for topic
topic_df = filtered_df[filtered_df["topic_label"] == topic]

# Automatically detect operational attributes
exclude_cols = ["business_id", "review", "topic", "topic_label", "sentiment", "sentiment_label"]
operational_attributes = [col for col in topic_df.columns if col not in exclude_cols]

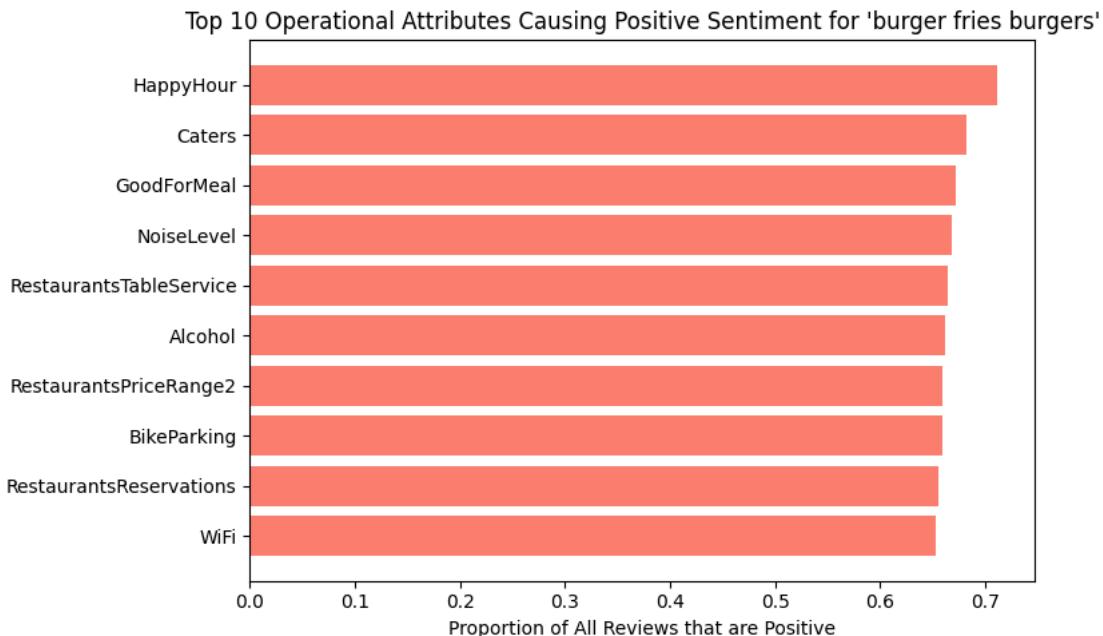
# Compute proportion of positive reviews for each attribute relative to all reviews
pos_attr_scores = {}

for attr in operational_attributes:
    if topic_df[attr].dtype == 'object':
        # Get all reviews for this attribute value
        counts = topic_df.groupby(attr)[["sentiment_label"]].value_counts(normalize=True)
        # Take the proportion of positive sentiment for the most common attribute value
        most_common_value = topic_df[attr].mode()[0]
        pos_score = counts.get((most_common_value, "Positive"), 0)
        pos_attr_scores[attr] = pos_score

# Sort descending and take top 5
top5_posAttrs = sorted(pos_attr_scores.items(), key=lambda x: x[1], reverse=True)[:10]

# Plot
attrs, scores = zip(*top5_posAttrs)

plt.figure(figsize=(8,5))
plt.barh(attrs, scores, color='salmon')
plt.xlabel("Proportion of All Reviews that are Positive")
plt.title(f"Top 10 Operational Attributes Causing Positive Sentiment for '{topic}'")
plt.gca().invert_yaxis() # highest at top
plt.tight_layout()
plt.show()
```



```

import matplotlib.pyplot as plt

def plot_top_negative_attributes(filtered_df, topic):
    # Define topic
    topic = topic

    # Filter for topic
    topic_df = filtered_df[filtered_df["topic_label"] == topic]

    # Automatically detect operational attributes
    exclude_cols = ["business_id", "review", "topic", "topic_label", "sentiment", "sentiment_label"]
    operational_attributes = [col for col in topic_df.columns if col not in exclude_cols]

    # Compute proportion of negative reviews for each attribute relative to all reviews
    neg_attr_scores = {}

    for attr in operational_attributes:
        if topic_df[attr].dtype == 'object':
            # Get all reviews for this attribute value
            counts = topic_df.groupby(attr)[["sentiment_label"]].value_counts(normalize=True)
            # Take the proportion of negative sentiment for the most common attribute value
            most_common_value = topic_df[attr].mode()[0]
            neg_score = counts.get((most_common_value, "Negative"), 0)
            neg_attr_scores[attr] = neg_score

    # Sort descending and take top 5
    top5_negAttrs = sorted(neg_attr_scores.items(), key=lambda x: x[1], reverse=True)[:10]

    # Plot
    attrs, scores = zip(*top5_negAttrs)

    plt.figure(figsize=(8,5))
    plt.barh(attrs, scores, color='salmon')
    plt.xlabel("Proportion of All Reviews that are Negative")
    plt.title(f"Top 10 Operational Attributes Causing Negative Sentiment for '{topic}'")
    plt.gca().invert_yaxis() # highest at top
    plt.tight_layout()
    plt.show()

```

```

import matplotlib.pyplot as plt

def plot_top_positive_attributes(filtered_df, topic):
    # Define topic
    topic = topic

    # Filter for topic
    topic_df = filtered_df[filtered_df["topic_label"] == topic]

    # Automatically detect operational attributes
    exclude_cols = ["business_id", "review", "topic", "topic_label", "sentiment", "sentiment_label"]
    operational_attributes = [col for col in topic_df.columns if col not in exclude_cols]

    # Compute proportion of positive reviews for each attribute relative to all reviews

```

```

pos_attr_scores = {}

for attr in operational_attributes:
    if topic_df[attr].dtype == 'object':
        # Get all reviews for this attribute value
        counts = topic_df.groupby(attr)[["sentiment_label"]].value_counts(normalize=True)
        # Take the proportion of positive sentiment for the most common attribute value
        most_common_value = topic_df[attr].mode()[0]
        pos_score = counts.get((most_common_value, "Positive"), 0)
        pos_attr_scores[attr] = pos_score

    # Sort descending and take top 5
top5_posAttrs = sorted(pos_attr_scores.items(), key=lambda x: x[1], reverse=True)[:10]

# Plot
attrs, scores = zip(*top5_posAttrs)

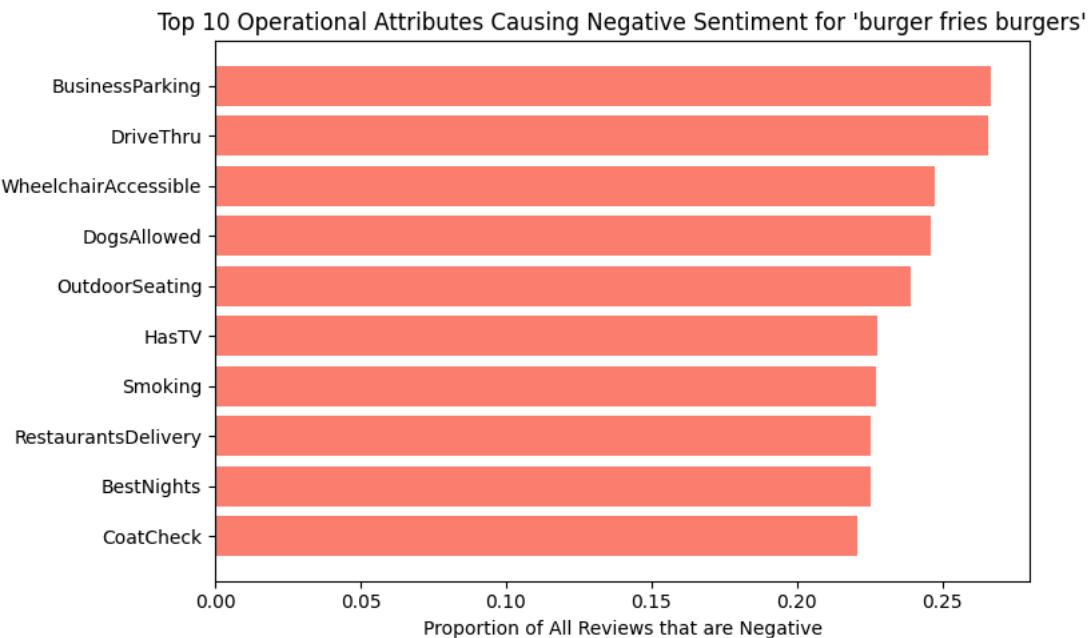
plt.figure(figsize=(8,5))
plt.barh(attrs, scores, color='salmon')
plt.xlabel("Proportion of All Reviews that are Positive")
plt.title(f"Top 10 Operational Attributes Causing Positive Sentiment for '{topic}'")
plt.gca().invert_yaxis() # highest at top
plt.tight_layout()
plt.show()

```

```

plot_top_negative_attributes(
    filtered_df,
    topic="burger fries burgers"
)

```

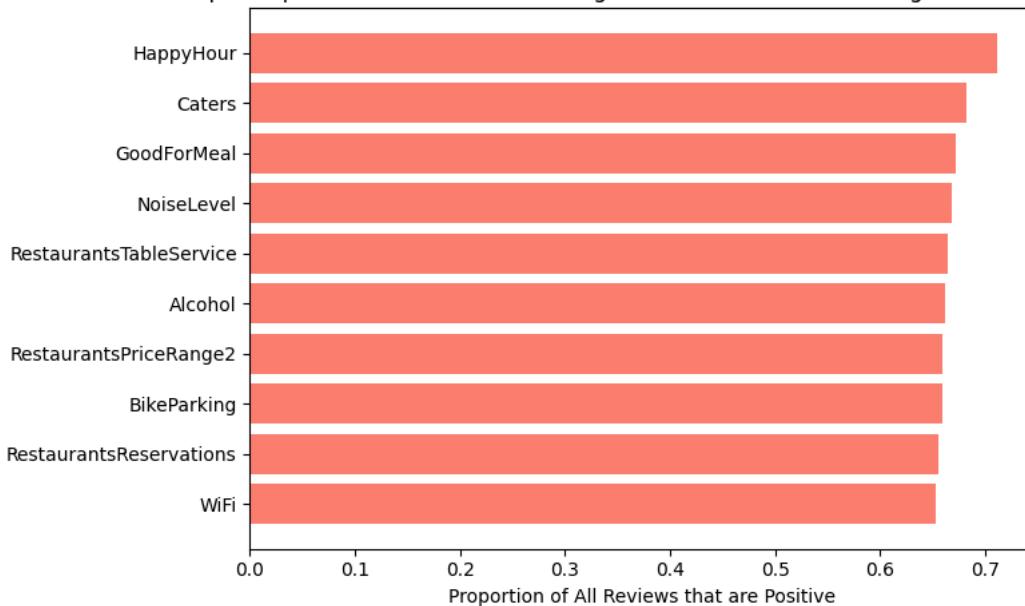


```

plot_top_positive_attributes(
    filtered_df,
    topic="burger fries burgers"
)

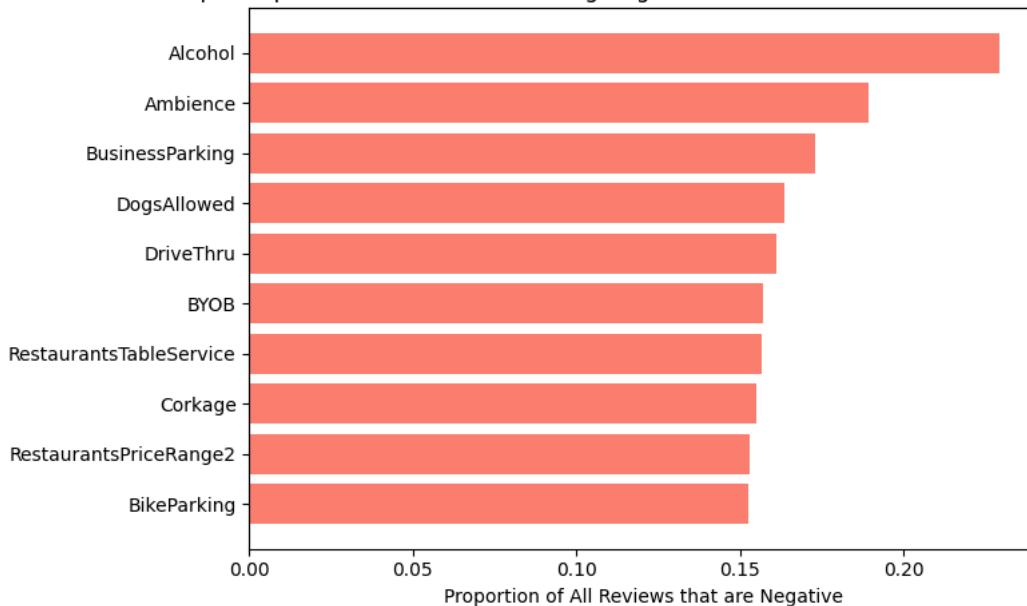
```

Top 10 Operational Attributes Causing Positive Sentiment for 'burger fries burgers'

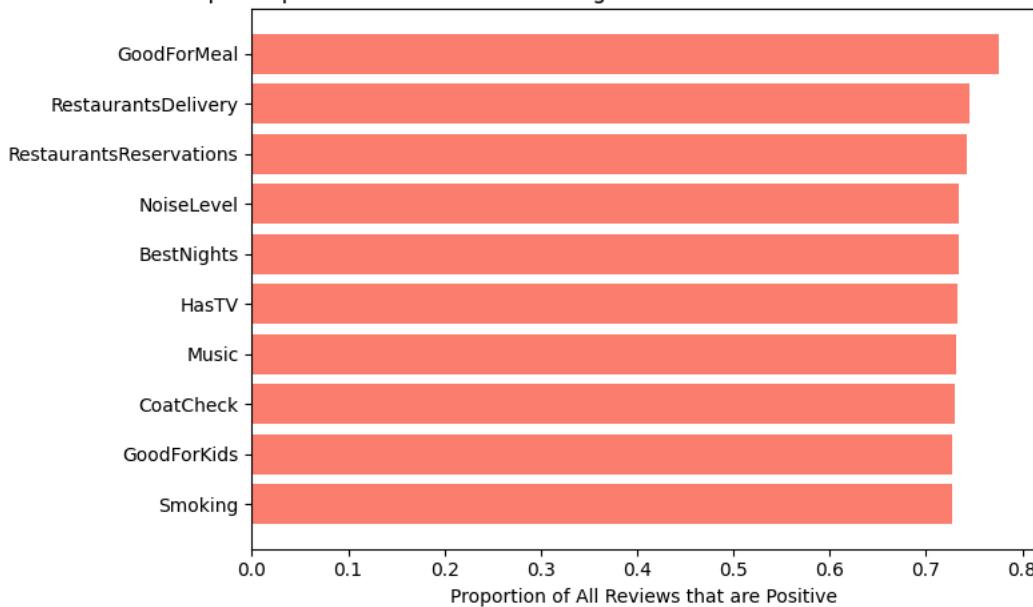


```
plot_top_negative_attributes(  
    filtered_df,  
    topic = "tacos mexican salsa"  
)  
  
plot_top_positive_attributes(  
    filtered_df,  
    topic = "tacos mexican salsa"  
)
```

Top 10 Operational Attributes Causing Negative Sentiment for 'tacos mexican salsa'



Top 10 Operational Attributes Causing Positive Sentiment for 'tacos mexican salsa'



Operational Attribute Impact Analysis

```

topic = "burger fries burgers"
attribute = "HappyHour"

topic_df = filtered_df[filtered_df["topic_label"] == topic]

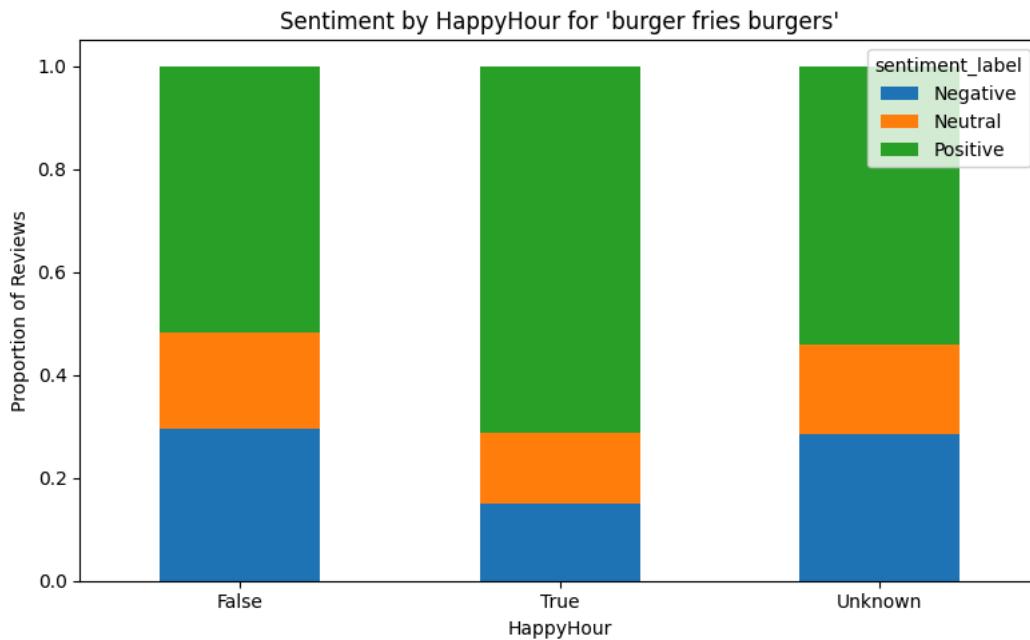
sentiment_dist = (
    topic_df
    .groupby([attribute, "sentiment_label"])
    .size()
    .unstack(fill_value=0)
)

sentiment_dist = sentiment_dist.div(sentiment_dist.sum(axis=1), axis=0)

sentiment_dist.plot(
    kind="bar",
    stacked=True,
    figsize=(8,5)
)

plt.title(f"Sentiment by {attribute} for '{topic}'")
plt.ylabel("Proportion of Reviews")
plt.xticks(rotation=0)
plt.tight_layout()
plt.show()

```



Pareto Analysis of Negative Sentiment Drivers

```

import matplotlib.pyplot as plt

#Count negative reviews per topic
neg_counts = (
    filtered_df[filtered_df["sentiment_label"] == "Negative"]
    ["topic_label"]
    .value_counts()
)

#Compute cumulative percentage
cum_pct = neg_counts.cumsum() / neg_counts.sum() * 100

#Keep only topics up to 80% cumulative impact
pareto_df = (
    cum_pct[cum_pct <= 80]
    .index
)

neg_counts_80 = neg_counts.loc[pareto_df]
cum_pct_80 = cum_pct.loc[pareto_df]

# Plot
plt.figure(figsize=(10,5))

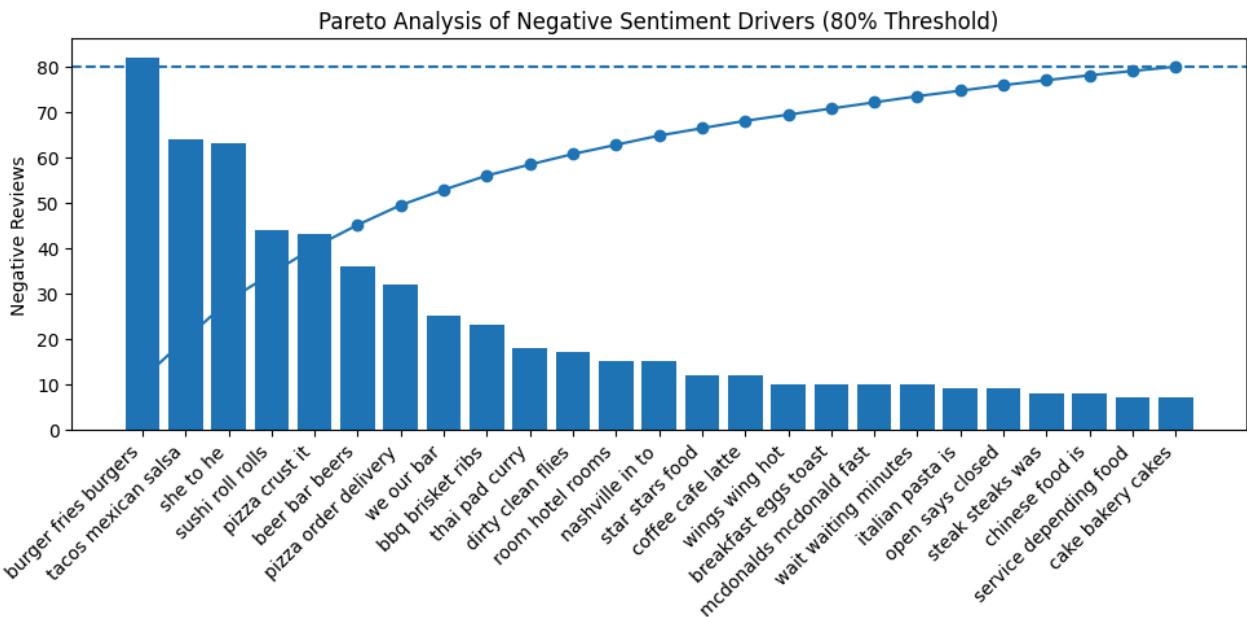
plt.bar(neg_counts_80.index, neg_counts_80.values)

```

```

        plt.plot(
            neg_counts_80.index,
            cum_pct_80.values,
            marker="o"
        )

        plt.axhline(80, linestyle="--")
        plt.ylabel("Negative Reviews")
        plt.title("Pareto Analysis of Negative Sentiment Drivers (80% Threshold)")
        plt.xticks(rotation=45, ha="right")
        plt.tight_layout()
        plt.show()
    
```



Heatmap Analysis of Operational Drivers

```

import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd

def plot_single_topic_attribute_heatmap(filtered_df, topic):
    #Filter for topic
    topic_df = filtered_df[filtered_df["topic_label"] == topic]

    #Exclude non-operational columns
    exclude_cols = ["business_id", "review", "topic", "topic_label", "sentiment", "sentiment_label"]
    attributes = [c for c in topic_df.columns if c not in exclude_cols]

    heatmap_data = []

    for attr in attributes:
        if topic_df[attr].dtype == "object":
            counts = (
                topic_df
                .groupby(attr)[["sentiment_label"]]
                .value_counts(normalize=True)
            )
            most_common_value = topic_df[attr].mode()[0]
            neg_score = counts.get((most_common_value, "Negative"), 0)

            heatmap_data.append({
                "attribute": attr,
                "negative_ratio": neg_score
            })
        else:
            heatmap_data.append({
                "attribute": attr,
                "negative_ratio": 0
            })

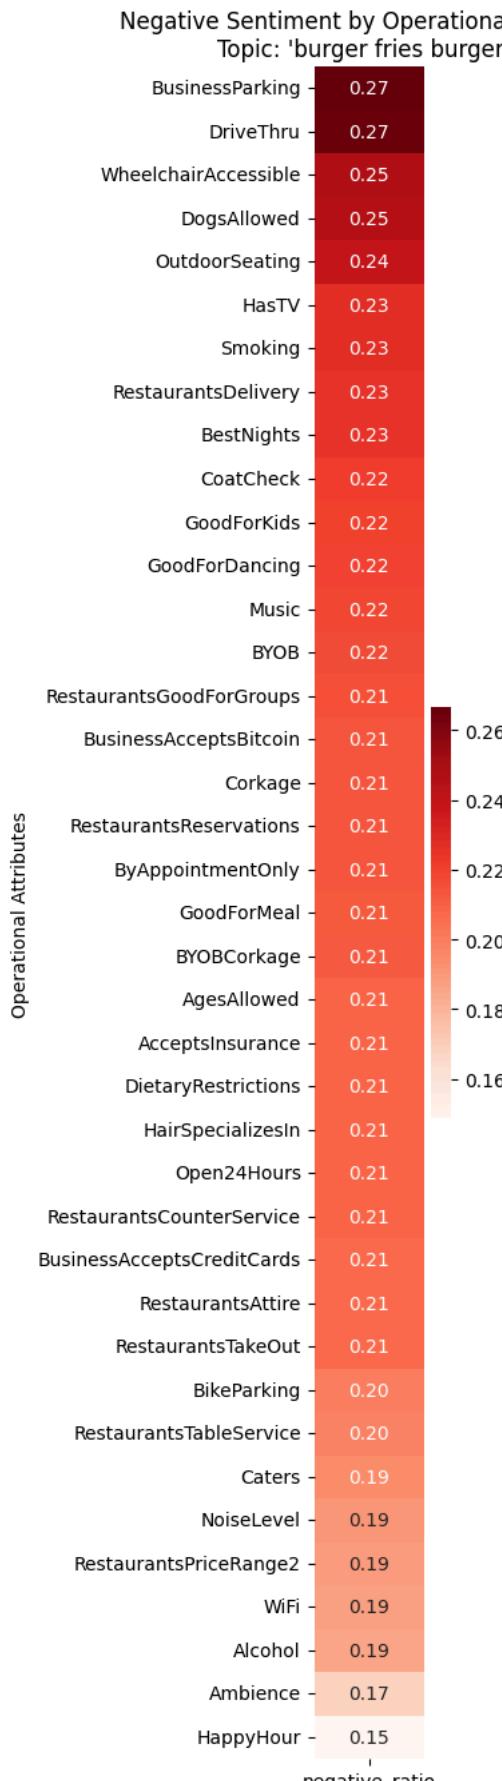
    heatmap_df = (
        pd.DataFrame(heatmap_data)
        .set_index("attribute")
        .sort_values("negative_ratio", ascending=False)
    )

    plt.figure(figsize=(4, max(6, len(heatmap_df) * 0.35)))
    sns.heatmap(
        heatmap_df,
        cmap="Reds",
        annot=True
    )

```

```
        annot=True,  
        fmt=".2f",  
        cbar=True  
    )  
  
    plt.title(f"Negative Sentiment by Operational Attribute\nTopic: '{topic}'")  
    plt.xlabel("")  
    plt.ylabel("Operational Attributes")  
    plt.tight_layout()  
    plt.show()
```

```
)  
  
plot_single_topic_attribute_heatmap(  
    filtered_df,  
    topic="burger fries burgers"
```



```
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd

def plot_single_topic_sentiment_heatmap(filtered_df, topic):
    #Filter for topic
    topic_df = filtered_df[filtered_df["topic_label"] == topic]

    #Exclude non-operational columns
```

```
exclude_cols = ["business_id", "review", "topic", "topic_label", "sentiment", "sentiment_label"]
attributes = [c for c in topic_df.columns if c not in exclude_cols]

heatmap_data = []

for attr in attributes:
    if topic_df[attr].dtype == "object":
        counts = (
            topic_df
            .groupby(attr)[["sentiment_label"]]
            .value_counts(normalize=True)
        )

        most_common_value = topic_df[attr].mode()[0]

        neg_ratio = counts.get((most_common_value, "Negative"), 0)
        pos_ratio = counts.get((most_common_value, "Positive"), 0)

        #Net sentiment score: positive - negative
        net_score = pos_ratio - neg_ratio

        heatmap_data.append({
            "attribute": attr,
            "net_sentiment": net_score
        })

heatmap_df = (
    pd.DataFrame(heatmap_data)
    .set_index("attribute")
    .sort_values("net_sentiment", ascending=True)
)

plt.figure(figsize=(4, max(6, len(heatmap_df) * 0.35)))
sns.heatmap(
    heatmap_df,
    cmap="RdYlGn",
    center=0,
    annot=True,
    fmt=".2f",
    cbar_kws={"label": "Net Sentiment (Positive - Negative)"}
```